



# Database Version Control with Alembic: Best Practices and Techniques

Pycon APAC 2025

2025-03-01

# Agenda

- What is alembic and why we need it?
- How to perform db migrations
- Microservices with shared DB

# What is alembic and why we need it?

---

# What is a Database Migration?

*It's a way to manage and track schema changes in your database.*

Examples:


- Create/rename/drop table.
- Add column.
- Relationship definitions.



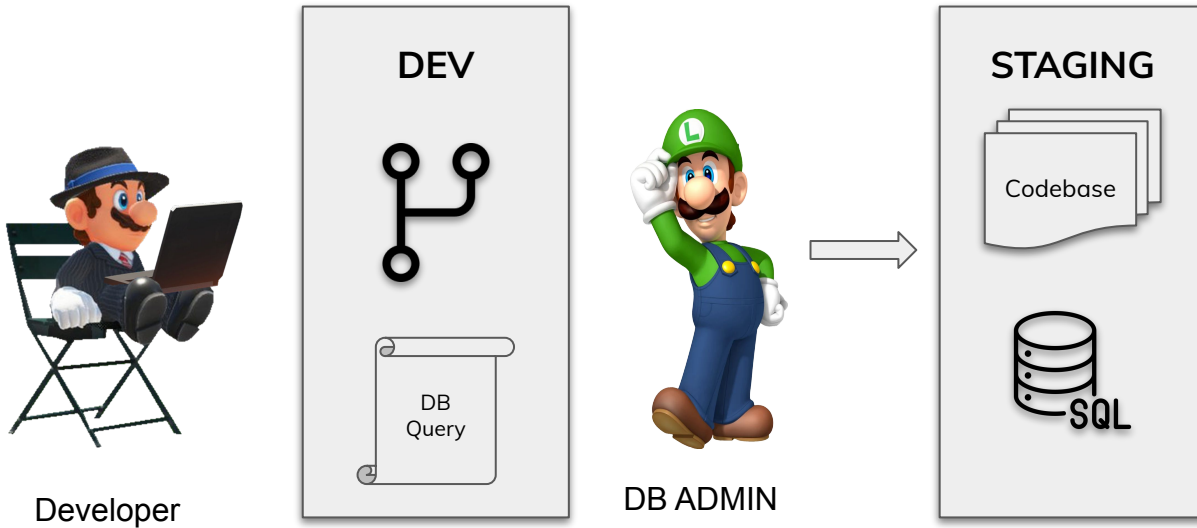
# Development workflow



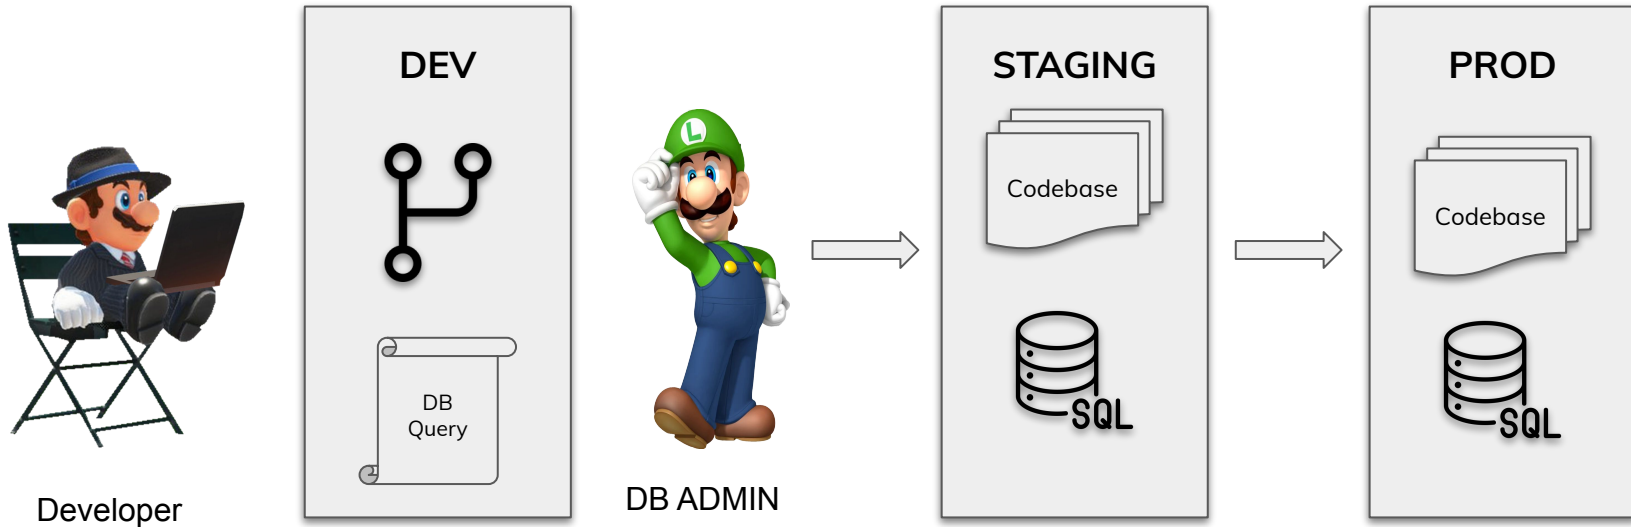
Developer



# Development workflow - Open PR

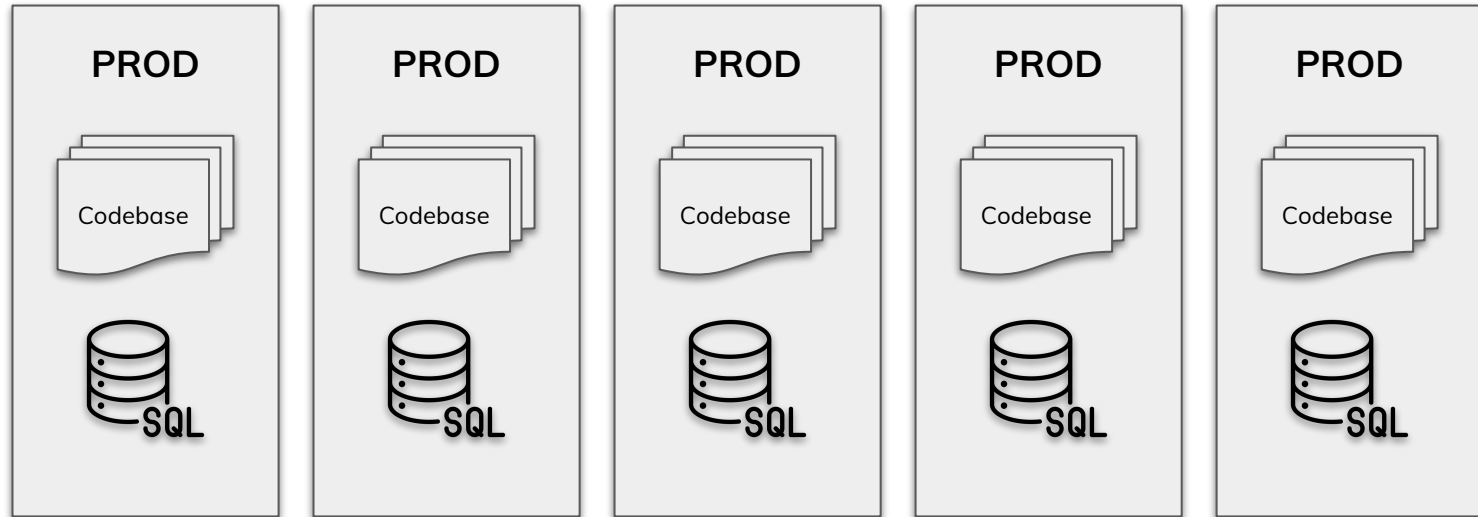


# Development workflow - Deploy



# Development workflow

What if your clients ask you to seclude their data?








# Enter Alembic!

*a lightweight migration tool for SQLAlchemy.*

- Track schema changes over time.
- Move between different revision.  
(upgrade/downgrade)
- Keep your database consistent across multiple environments.





## pip install alembic

After running:

```
~/pycon-apac$ alembic init alembic
```

We get the following structure:

```

📁 alembic/
├── env.py      # Main configuration file
├── script.py.mako # Template for new migrations
├── 📁 versions/  # Stores migration scripts
├── README      # Basic documentation
└── 📄 alembic.ini # Global settings

```

# Object-Relational Mapping (ORM)

- ORM abstracts SQL → No longer raw queries.
- Table becomes Python classes.
- SQLAlchemy manages connections, queries, transactions.

## Example:

```

model.py > ...
1  from sqlalchemy import Column, String, Float, DateTime
2  from sqlalchemy.ext.declarative import declarative_base
3  from sqlalchemy.orm import validates
4
5  Base = declarative_base()
6
7  class Prices(Base):
8      __tablename__ = 'prices'
9
10     date = Column(DateTime, primary_key=True)
11     equity_id = Column(String(32), primary_key=True)
12     price = Column(Float, nullable=False)
13     volume = Column(Float, nullable=False)

```

# How to perform db migrations

---

## Create first revision

```
roberto-landi@l-100074-sh19:~/pycon-apac$ alembic revision -m "first_revision"  
  Generating /home/roberto-landi/pycon-apac/alembic/versions/463161aec4ec_first_revision.py ...  done  
roberto-landi@l-100074-sh19:~/pycon-apac$ █
```

# Migration file structure

```
alembic > versions > 463161aec4ec_first_revision.py > ...
1  """first_revision
2
3     Revision ID: 463161aec4ec
4     Revises:
5     Create Date: 2025-02-04 17:07:20.132441
6
7     """
8  from typing import Sequence, Union
9
10 from alembic import op
11 import sqlalchemy as sa
12
13
14 # revision identifiers, used by Alembic.
15 revision: str = '463161aec4ec'
16 down_revision: Union[str, None] = None
17
18
19 def upgrade() -> None:
20     pass
21
22
23 def downgrade() -> None:
24     pass
25
```



# Write first revision

```

18 def upgrade() -> None:
19     op.create_table('prices',
20         sa.Column('date', sa.DateTime(), nullable=False),
21         sa.Column('equity_id', sa.String(length=32), nullable=False),
22         sa.Column('price', sa.Float(), nullable=False),
23         sa.Column('volume', sa.Float(), nullable=False),
24         sa.PrimaryKeyConstraint('date', 'equity_id')
25     )
26
27 def downgrade() -> None:
28     op.drop_table('prices')
29

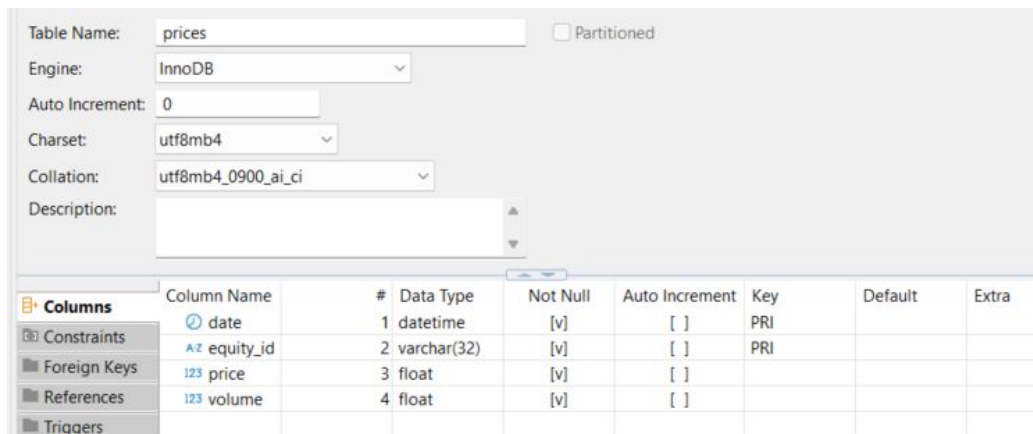

```

# Database view

 workshop/

└─ alembic\_version

└─ prices



# Add data quality checks

```

4  base = declarative_base()
5
6  class Prices(Base):
7      __tablename__ = 'prices'
8      __table_args__ = (UniqueConstraint('date', 'equity_id', 'market', name='unique_date_equity_id_market'),)
9
10     date = Column(DateTime, primary_key=True)
11     equity_id = Column(String(32), primary_key=True)
12     market = Column(String(32), nullable=True)
13     price = Column(Float, nullable=False)
14     volume = Column(Float, nullable=False)
15

```

Ctrl+L to chat, Ctrl+K to generate

# Alembic revision –autogenerate

Alembic can **automatically detect** schema changes in your SQLAlchemy models and generate migration scripts for you.

```
INFO [alembic.runtime.migration] Running downgrade 6587ee11a882 -> 405161a2c4ec; data_quality
• (venv) roberto-landi@l-100074-sh19:~/pycon-apac$ alembic revision --autogenerate -m "quality_layer"
INFO [alembic.runtime.migration] Context impl MySQLImpl.
INFO [alembic.runtime.migration] Will assume non-transactional DDL.
INFO [alembic.autogenerate.compare] Detected added column 'prices.market'
INFO [alembic.autogenerate.compare] Detected added unique constraint 'unique_date_equity_id_market' on '('date', 'equity_id', 'market')'
Generating /home/roberto-landi/pycon-apac/alembic/versions/9b84540a7c1c_quality_layer.py ... done
○ (venv) roberto-landi@l-100074-sh19:~/pycon-apac$
```

# Autogenerated migration

```

alembic > versions > 9b84540a7c1c_quality_layer.py > ...
12
13
14 # revision identifiers, used by Alembic.
15 revision: str = '9b84540a7c1c'
16 down_revision: Union[str, None] = '463161aec4ec'
17 branch_labels: Union[str, Sequence[str], None] = None
18 depends_on: Union[str, Sequence[str], None] = None
19
20
21 def upgrade() -> None:
22     # ### commands auto generated by Alembic - please adjust! ###
23     op.add_column('prices', sa.Column('market', sa.String(length=32), nullable=True))
24     op.create_unique_constraint('unique_date_equity_id_market', 'prices', ['date', 'equity_id', 'market'])
25     # ### end Alembic commands ###
26
27
28 def downgrade() -> None:
29     # ### commands auto generated by Alembic - please adjust! ###
30     op.drop_constraint('unique_date_equity_id_market', 'prices', type_='unique')
31     op.drop_column('prices', 'market')
32     # ### end Alembic commands ###
33

```

# Alembic upgrade head

```
(venv) roberto-landi@l-100074-sh19:~/pycon-apac$ alembic upgrade head
INFO [alembic.runtime.migration] Context impl MySQLImpl.
INFO [alembic.runtime.migration] Will assume non-transactional DDL.
INFO [alembic.runtime.migration] Running upgrade 463161aec4ec -> 9b84540a7c1c, quality_layer
(venv) roberto-landi@l-100074-sh19:~/pycon-apac$
```

Name:

Owner: [workshop.prices](#)

Type:

Check expression:

Name	Column
date	<a href="#">date</a>
<a href="#">A-Z</a> equity_id	<a href="#">A-Z</a> <a href="#">equity_id</a>
<a href="#">A-Z</a> market	<a href="#">A-Z</a> <a href="#">market</a>

Properties		Data	ER Diagram
alembic_version		Enter a SQL expression to filter	
Grid		<a href="#">A-Z</a> version_num	
		1	9b84540a7c1c
Text			

# Renaming columns

```

3
4 Base = declarative_base()
5
6 class Prices(Base):
7     __tablename__ = 'prices'
8     __table_args__ = (UniqueConstraint('date', 'equity_id', 'mic', name='unique_date_equity_id_market'),)
9
10    date = Column(DateTime, primary_key=True)
11    equity_id = Column(String(32), primary_key=True)
12    mic = Column(String(32), nullable=True)
13    #market = Column(String(32), nullable=True)
14    price = Column(Float, nullable=False)
15    volume = Column(Float, nullable=False)
16

```



# Alembic revision –autogenerate

```

• (venv) roberto-landi@1-100074-sh19:~/pycon-apac$ alembic revision --autogenerate -m "rename mic"
INFO [alembic.runtime.migration] Context impl MySQLImpl.
INFO [alembic.runtime.migration] Will assume non-transactional DDL.
INFO [alembic.autogenerate.compare] Detected added column 'prices.mic'
INFO [alembic.autogenerate.compare] Detected changed unique constraint 'unique_date_equity_id_market' on 'prices': expression ('date', 'equity_id', 'market') to ('date', 'equity_id', 'mic')
INFO [alembic.autogenerate.compare] Detected removed column 'prices.market'
Generating /home/roberto-landi/pycon-apac/alembic/versions/4199f6c54dce_rename_mic.py ... done
○ (venv) roberto-landi@1-100074-sh19:~/pycon-apac$

```



# Alembic upgrade head

```

generating /home/roberto-landi/pycon-apac/alembic/versions/4199f6c54dce_rename_mic.py ...
• (venv) roberto-landi@l-100074-sh19:~/pycon-apac$ alembic upgrade head
INFO [alembic.runtime.migration] Context impl MySQLImpl.
INFO [alembic.runtime.migration] Will assume non-transactional DDL.
INFO [alembic.runtime.migration] Running upgrade 9b84540a7c1c -> 4199f6c54dce, rename mic
○ (venv) roberto-landi@l-100074-sh19:~/pycon-apac$

```

# Alembic upgrade head



```
generating /home/roberto-landi@l-100
• (venv) roberto-landi@l-100
INFO [alembic.runtime.mig
INFO [alembic.runtime.mig
INFO [alembic.runtime.mig
○ (venv) roberto-landi@l-100
```

```
654dce_rename_mic.py ...
ad
.
199f6c54dce, rename mic
```

# Autogenerated migration

```

15 revision: str = '4199f6c54dce'
16 down_revision: Union[str, None] = '9b84540a7c1c'
17 branch_labels: Union[str, Sequence[str], None] = None
18 depends_on: Union[str, Sequence[str], None] = None
19
20
21 def upgrade() -> None:
22     # ### commands auto generated by Alembic - please adjust! ###
23     op.add_column('prices', sa.Column('mic', sa.String(length=32), nullable=True))
24     op.drop_constraint('unique_date_equity_id_market', 'prices', type_='unique')
25     op.create_unique_constraint('unique_date_equity_id_market', 'prices', ['date', 'equity_id', 'mic'])
26     op.drop_column('prices', 'market')
27     # ### end Alembic commands ###
28
29
30 def downgrade() -> None:
31     # ### commands auto generated by Alembic - please adjust! ###
32     op.add_column('prices', sa.Column('market', mysql.VARCHAR(length=32), nullable=True))
33     op.drop_constraint('unique_date_equity_id_market', 'prices', type_='unique')
34     op.create_unique_constraint('unique_date_equity_id_market', 'prices', ['date', 'equity_id', 'market'])
35     op.drop_column('prices', 'mic')
36     # ### end Alembic commands ###

```

# What does Autogenerate Detect ?

## Detects :

- Table/columns additions and removals.
- Column type changes.
- Simple constraints.


## Does NOT Detect :

- Change of column/table names.
- Complex constraints.
- Data migrations.

# Microservices with shared DB

---

# Microservices with shared DB



## Microservices with shared DB - What Goes Wrong?

- How do I keep revision versions aligned?
- How do I prevent table deletions and data loss?

# Microservices with shared DB - env.py

```

62 def run_migrations_online() -> None:
63     """Run migrations in 'online' mode.
64
65     In this scenario we need to create an Engine
66     and associate a connection with the context.
67
68     """
69     connectable = engine_from_config(
70         config.get_section(config.config_ini_section, {}),
71         prefix="sqlalchemy.",
72         poolclass=pool.NullPool,
73     )
74
75     with connectable.connect() as connection:
76         context.configure(
77             connection=connection,
78             target_metadata=target_metadata,
79             version_table="pycon_versions",
80             include_object=include_object,
81         )
82
83     with context.begin_transaction():
84         context.run_migrations()
85


```




# Microservices with shared DB - env.py

```
def include_object(object, name, type_, reflected, compare_to):
    if type_ == "table" and name.startswith("pycon"):
        return True
    elif type_ == "table":
        return False
    else:
        return True
```

```
5
6 class Prices(Base):
7     __tablename__ = 'pycon_prices'
8     __table_args__ = (UniqueConstraint('date'),)
9     Ctrl+L to chat, Ctrl+K to generate
10    date = Column(DateTime, primary_key=True)
```



# Microservices with shared DB - Solution



## Key takeaways

- Database migrations ensure schema changes are applied incrementally and consistently.
- Alembic provides version-controlled database migrations, like Git for databases.
- Autogenerate is helpful, but always review migrations before applying them.
- Alembic is best for schema changes, not data migrations.
- Handling multiple projects with a shared DB requires custom migration strategies.

# Salamat Po!



Roberto Landi, Head of data engineering.



roberto.landi@axyon.ai



<https://axyon.ai>

